

## experience

### Hugging Face

Research Engineer & Team Lead, Research Reproduction January 2023 - present

- Leading large-scale language modeling efforts within the new Research Reproduction team

Research Engineer, Science team

July 2022 - December 2023

- Developed pipelines and tools for zero-shot evaluation of large language models in the Evaluate library and in Evaluation on the Hub ([blog](#))
- Wrote pipelines on top of DeepSpeed to evaluate the OPT series of language models on the WinoBias task and conducted research on mitigations through prompting at scale ([talk](#) on partial work)
- Created novel large evaluation datasets for temporal evaluation of language models by scraping websites for up-to-date data on real-world events as part of the [Online Language Modeling](#) project

### The AI Index, Stanford Human-Centered Artificial Intelligence (HAI)

Affiliated Researcher

November 2021 - March 2022, October 2022 - present

- Conducted meta-analysis on research literature for new chapter surveying technical metrics on bias and fairness, ethics, and conference publication trends for the 2022 [AI Index Report](#), 100k+ downloads annually and read by policymakers in the US and UK

### Cohere

Research Engineer, Safety Team Lead

January 2020 - March 2022

- Founding engineer; designed and implemented core model training framework in Tensorflow used to train and evaluate multi-billion parameter Transformer language models on TPUs; later rewritten in JAX
- Designed and implemented [safety system](#) using language models for dataset curation & filtration of toxic and undesirable web data resulting in models which minimize safety taxes
- Led safety research team responsible for building technical measurement and mitigation systems
- Wrote public-facing documentation communicating best practices around responsible use and collaborated with product managers to write best practices around product safety and policy

### Recurse Center

Fellowship Recipient

Fall 2019

- Researched misinformation and the integrity of our information ecosystem on Twitter. Built dashboards on Canadian elections & disinformation in Hong Kong deployed as Flask applications

### Dessa (acquired by Square Inc.)

Machine Learning Engineer

January - December 2019

- Designed, built, and shipped production machine learning systems for Fortune 100 enterprises at multi-terabyte scale across a variety of industries using Python, Tensorflow, Spark, Docker, and SQL
- Key projects: credit card fraud detection for top-3 U.S. creditor, social media diversity & inclusion brand sentiment analysis tool for [Diversio](#), consulted on credit card fraud system for Canadian bank

### Sidewalk Labs & Waterfront Toronto

Research Fellow

March - December 2018

- Selected as 1 of 12 Fellows as part of Sidewalk Toronto's mission to build a next-generation smart city
- Consulting on technology and policy with the goal of improving urban life for all. Report: [link](#)

Bell

Data Scientist, AI Labs team

June 2016 - December 2018

- Built, trained and deployed deep learning classifier for call center prediction; 20% lift over baseline
- Regression & tree-based modeling for churn, pricing sensitivity, iconic device launches, ad targeting
- Integrated models into production environment creating 100M+ scores weekly

## research

### Generative Language Models and Automated Influence Operations (2023)

Workshop participant for preprint published in 2023

### Evaluate & Evaluation on the Hub: Better Best Practices for Data and Model Measurements (2022)

Leandro von Werra, Lewis Tunstall, Abhishek Thakur, ... **Helen Ngo**, ... Thomas Wolf, Douwe Kiela  
Accepted to EMNLP 2022 demo track and selected for Best Demo Award

### No News is Good News: A Critique of the One Billion Word Benchmark (2021)

**Helen Ngo**, João GM Araújo, Jeffrey Hui, Nicholas Frosst

Accepted to the Data-Centric Workshop @ NeurIPS 2021. Selected as a lightning talk.

### Mitigating Harm in Language Models with Conditional-Likelihood Filtration (2021)

**Helen Ngo**, Cooper Raterink, João GM Araújo, Ivan Zhang, Carol Chen, Nicholas Frosst

### Predicting Twitter Engagement with Deep Language Models (2020)

M Volkovs, Z Cheng, M Ravaut, H Yang, K Shen, JP Zhou, A Wong, S Zuberi, I Zhang, N Frosst, **H Ngo**, C Chen, B Venkitesh, S Gou, A.N. Gomez

Proceedings of the Recommender Systems Challenge 2020

## awards, writing, miscellaneous

- Writer-in-Residence @ The University of Western Ontario (2015-2016), selected from 400 applicants to be first STEM student named to Writer-in-Residence post
- **Book:** [Reimagining ChinaTown: Speculative Stories from Toronto's Chinatown](#) (Mawenzi House 2023)
- **Community service:** Recurse Center alumni interview (2022 - present), Toronto Women's Data Group organizer (2017 - 2020)
- **Projects:** [@whalefakes](#) (finetuned GPT-2 Twitter bot), [Wolfram Beta](#) (Transformers for mathematics)
- **Awards:** OpenAI Scholar 2020 (declined), Recurse Center fellowship 2019, ACM Recsys Challenge 2019 (top 3% of teams worldwide), VentureBeat Women in AI Award finalist 2019, PyData NYC Scholarship 2019 (declined), Dockercon Scholarship 2018

## education

The University of Western Ontario

2016

Honors Specialization, Mathematics (B.Sc), graduated on Dean's Honor List

Minor in Writing Studies

**Awards:** Alfred Poynt Poetry Award, Marie Smibert Writing Prize, Barbara Newbegin Award in Mathematics